

Preview

- Process Termination
- Zombie Process
- wait() and waitpid() System Call
- Orphan Process
- exec System Calls
 - execl()
 - execv()
 - execl_e()
 - execve()
 - execlp()
 - execvp()

COSC350 System Software, Fall 2020
Dr. Sang-Eon Park

1

Process Termination

- No matter how a process terminated normally or abnormally, kernel execute a code closes all the open descriptors, release the memory used and so on.
- When a process terminated, the parents can obtain child's status from either the wait() or the waitpid() system call.
- If a parent terminates before the child, **systemd**(init in Unix) process becomes the parent process of any process whose parent terminated.

COSC350 System Software, Fall 2020
Dr. Sang-Eon Park

2

Process Termination

- When a parent process terminate, the parent process ID of the surviving process is changed to be 1 (Guaranteed every process has a parent)
 - Process ID = 0 : scheduler process
 - Process ID = 1 : systemd (init in Unix) process
- If a child terminate before parent, the kernel save a child's information (ID, termination status, CPU time) for the parent process termination.

COSC350 System Software, Fall 2020
Dr. Sang-Eon Park

3

Process Termination

- When a child process finishes execution, it will have an exit status to report to its parent process.
- Because of this last little bit of information, the process will remain in the operating system's process table as a **zombie process**.
- A **zombie process** will not to be scheduled for further execution, but that it cannot be completely removed
- `ps -el | grep 'Z'` can prints the status of zombie.

COSC350 System Software, Fall 2020
Dr. Sang-Eon Park

4

Process Termination

- When a child exits, its parent is supposed to use the "**wait**" system call and collect the child process's exit information.
- The subprocess exists as a zombie process until this happens.
- However, if the parent process isn't programmed properly or has a bug and never calls "**wait**," the zombie process remains, eternally waiting for its information to be collected by its parent.

COSC350 System Software, Fall 2020
Dr. Sang-Eon Park

5

Process Termination

```
/* zombie.c: create a zombie process */
#include <stdio.h>
#include <string.h>
#include <unistd.h>
#include <sys/types.h>

int main(void)
{
    int pid;
    pid=fork();
    /* parent sleep 100 sec */
    if(pid>0)
    {
        sleep(100);
    }
    /* a child terminate parent don't have time to save a child info */
    /* child remains as a zombie */
    else
    {
        exit(0);
    }
}
```

COSC350 System Software, Fall 2020
Dr. Sang-Eon Park

6

Process Termination

```
/* zombie0.c: create a zombie process */
#include <stdio.h>
#include <string.h>
#include <unistd.h>
#include <sys/types.h>
int main(void)
{
    int pid, pid1;
    pid=fork(); //create the first child
    if(pid>0)
    {
        pid1=fork(); //create the second child
        if (pid1 >0) //parent runs forever
        {
            while (1)
            {
                ;
            }
        }
        else // for second child
        {
            _exit(0);
        }
    }
    else //for first child
    {
        _exit(0);
    }
}
```

Process Termination

- ❑ The exit status of a child will be used for parent process termination.
- ❑ When a child exits, the parent process will receive a SIGCHLD signal to indicate that one of its children has finished executing; the parent process will typically call the wait() system call at this point.
- ❑ That call will provide the parent with the child's exit status, and will cause the child to be reaped, or removed from the process table.

COSC350 System Software, Fall 2020
Dr. Sang-Eon Park

8

wait and waitpid() System Call

- ❑ When a process terminate either normally or abnormally, the kernel sent a signal (SIGCHD) to a parent.
- ❑ A parent can ignore the signal or call a function (wait or waitpid) to take care the signal.

```
#include <sys/wait.h>
pid_t wait (int *status);
pid_t waitpid(pid_t pid, int *status, int option);
```

COSC350 System Software, Fall 2020
Dr. Sang-Eon Park

9

wait and waitpid() System Call

- ❑ The execution of **wait()** could have two possible situations.
 - If there are at least one child processes, the caller will be blocked until one of its child processes exits.
 - If there is no child process running, then this **wait()** has no effect at all.
- ❑ The status is the pointer where terminated process's status is saved.

COSC350 System Software, Fall 2020
Dr. Sang-Eon Park

10

Process Termination

```
/* zombie01.c: create a zombie process */
#include <stdio.h>
#include <string.h>
#include <unistd.h>
#include <sys/types.h>
int main(void)
{
    int pid, pid1;
    pid=fork(); //create the first child
    if(pid>0)
    {
        pid1=fork(); //create the second child
        if (pid1 >0) //parent runs forever
        {
            wait();
            while (1)
            {
                ;
            }
        }
        else // for second child
        {
            _exit(0);
        }
    }
    else //for first child
    {
        _exit(0);
    }
}
```

wait and waitpid() System Call

- ❑ By using macros in <sys/wait.h>, we can check a terminated process's status.
- ❑ The status field will be filled in by wait or waitpid function.

```
#include <sys/wait.h>
int WIFEXITED(int status);
    if child process terminate normally, return true

int WIFSIGNALED(int status);
    if child process terminate abnormally, return true
```

COSC350 System Software, Fall 2020
Dr. Sang-Eon Park

12

```

/*exit.c demonstrate wait() system call */
#include <stdio.h>
#include <string.h>
#include <sys/wait.h>
#include <unistd.h>
#include <sys/types.h>
#define MAX_COUNT 1000
#define BUF_SIZE 100

void main(void)
{
    pid_t ppid, pid, p1;
    int i;
    char buf[BUF_SIZE];
    printf("Parent is about to fork ***\n");
    if ((ppid = fork()) < 0)
    {
        printf("Failed to fork process 1\n");
        wait(1);
    }
    else if (ppid == 0) /* child's part */
    {
        p1 = getpid();
        for (i = 1; i < MAX_COUNT; i++)
        {
            sprintf(buf, "child %d is running\n", p1);
            write(1, buf, strlen(buf));
        }
        _exit(0);
    }
    else /* parent's part */
    {
        printf("Parent enters waiting status .....n");
        wait(1, buf, strlen(buf));
        pid = wait(&status); /*wait for child finish it's job */
        if (WEXITSTATUS(status))
            printf("A child process terminate normally\n");
        else
            printf("A child process terminate abnormally\n");
        sprintf(buf, "Parent detects process %d was done ***\n", pid);
        write(1, buf, strlen(buf));
        printf("Parent exits ***\n");
        wait(0);
    }
}

```

Orphan Process

- An **Orphan Process** is nearly the same thing which we see in real world.
- Orphan means someone whose parents are dead.
- The same way this is a process, whose parents are dead, that means parents are either terminated, killed or exited but the child process is still alive.

COSC350 System Software, Fall 2020
Dr. Sang-Eon Park

14

```

/* orphan.c shows example orphan call */
#include <stdio.h>
#include <string.h>
#include <sys/types.h>

int main (int argc, const char *argv[])
{
    pid_t pid;
    char buf[100];
    if ((pid = fork()) < 0) /* create child */
    {
        printf("fork error\n");
        exit(1);
    }
    if (pid == 0) /* child process */
    {
        while (1)
        {
            sprintf(buf, "**** Child %d is running. Its parent is %d .....n", getpid(), getppid());
            write(1, buf, strlen(buf));
            sleep(1);
            if (getppid() == 1)
                printf("My parent is passed away. Now, I am a orphan with step parent %d\n", getppid());
        }
    }
    else /* parent process */
    {
        while (1)
        {
            sprintf(buf, "**** %d am still alive\n", getpid());
            write(1, buf, strlen(buf));
            sleep(1);
        }
    }
    exit(0);
}

```

exec System Call

- By using exec system call, a child process can execute another program.
- Once a process call a exec system call, that process is completely replaced by the new program.
- The new program starts executing at its main function. The main function might need arguments.
- The process ID does not change across an exec system call, since it is not created.
- The content of text, data, heap and stack segment will be replaced by new program.

COSC350 System Software, Fall 2020
Dr. Sang-Eon Park

15

exec System Call

```

#include <unistd.h>
int execl(const char *path, const char *arg0, ... /*, (char *)0 */);
int execv(const char *path, char *const argv[]):
int execl(const char *path, const char *arg0, ... /*, (char *)0, char *const envp[] */);
int execve(const char *path, char *const argv[], char *const envp[]):
int execlp(const char *file, const char *arg0, ... /*, (char *)0 */);
int execvp(const char *file, char *const argv[]):

```

Return -1 on error, no return on success

COSC350 System Software, Fall 2020
Dr. Sang-Eon Park

17

exec System Call

- Six system call can be recognized by
 - Argument list or Argument vector
 - File name or path name
 - With or without environment

COSC350 System Software, Fall 2020
Dr. Sang-Eon Park

18

exec System Call

- ❑ If *filename* contains a slash, it is consider as a pathname.
- ❑ Otherwise, the executable file is searched for the directory specified PATH environment variable.
- ❑ If a file name find out but not executable, then it is consider as shell script and tries to invoke /bin/sh.
- ❑ With `execl` and `execve`, environment variable can be passed to the function.

COSC350 System Software, Fall 2020
Dr. Sang-Eon Park

19

exec System Call

- ❑ Normally, a process allows its environment to be propagated to its children.
- ❑ But some cases, a process need to specify a certain environment for a child.
Ex) the login program need create different environment for each user's login.

COSC350 System Software, Fall 2020
Dr. Sang-Eon Park

20

```
/* exexec.c shows execv system call */
#include <stdio.h>
#include <string.h>
#include <sys/types.h>

int main (int argc, const char *argv[])
{
    pid_t pid;
    int mult = 1, i;
    if (argc == 1)
    {
        printf("argument error \n");
        exit (1);
    }
    if ( (pid = fork()) < 0) /* create child */
    {
        printf("fork error\n");
        exit (1);
    }
    else if (pid == 0)
    {
        /* a child execute different program */
        if (execv ("/home/separk/Lecture/cosc350/example/ch4/summation", argv) < 0)
        {
            printf ("execl ERROR");
            exit (1);
        }
    }

    for (i = 1; i < argc; i++)
        mult *= atoi(argv[i]);
    printf ("The multiplication of arguments is %d \n", mult);
    exit (0);
}
```

COSC350 System Software, Fall 2020
Dr. Sang-Eon Park

21